

# Celery

## A Distributed Task Queue



Oskar Hollmann

User Technologies

26.10.2016

# MOTIVATION

- ▶ time demanding tasks are a pain in web apps
  - ▶ HTTP request can easily time out
  - ▶ it's not acceptable to block the client for too long
  - ▶ client may not care about the result
- ▶ either we return an URL to the client who polls it later to get the result
- ▶ or we push it through web sockets
- ▶ how can we achieve that?

# MOTIVATION

## Node.js to the rescue!



# MOTIVATION

*Node.js is cool, but...*

- ▶ Why use Node.js and struggle with async code when async operations are seldom needed in a web app?
- ▶ Async code is not enough – we might need to distribute the tasks, run them periodically, ...
- ▶ Node.js programmers are a rare commodity.
- ▶ What we actually need is **an asynchronous task queue**.
- ▶ Examples: RabbitMQ, JMS, Celery, ...



# USE CASES OF DISTRIBUTED TASK QUEUES

1. Non-blocking task execution.
2. Task execution with failure recovery.
3. Concurrent task execution for single-threaded apps.
4. Distribute task to other machines.
5. Handle complex task workflows with dependencies.
6. Periodic tasks.

# TASK QUEUE IN WEB APP

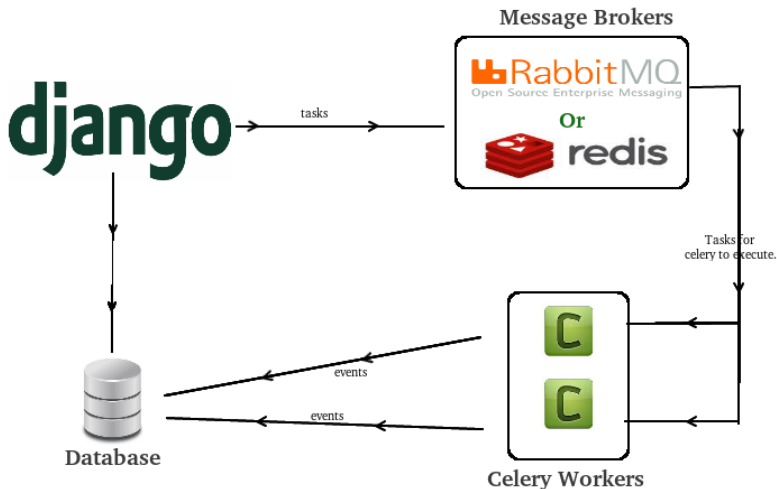


Figure: Source [en.proft.me](http://en.proft.me)

# CELERY IS

- ▶ distributed task queue written in Python
- ▶ bindings for: PHP, Ruby, NodeJS and more
- ▶ different message broker transports: Redis, RabbitMQ, MongoDB and more
- ▶ arbitrary number of queues and workers

# WORKERS AND QUEUES

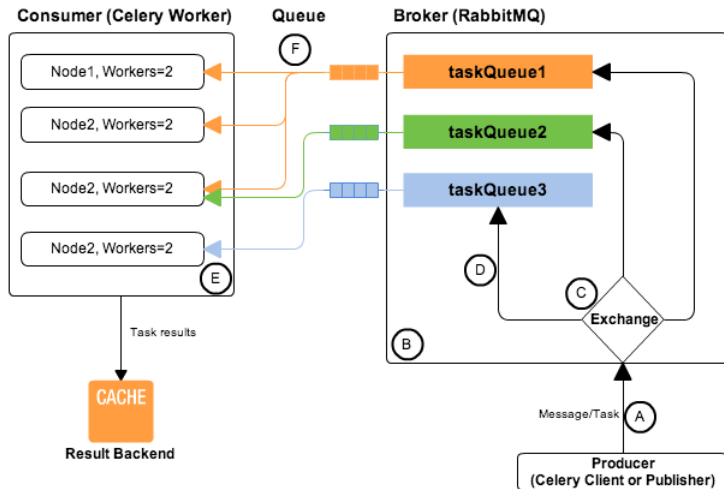


Figure: Source [abhishek-tiwari.com](http://abhishek-tiwari.com)



# MINIMAL EXAMPLE – CREATE CELERY APP

```
1  from celery import Celery
2
3  app = Celery(
4      'tasks',
5      broker='redis://localhost:6379/1',
6      backend='redis://',
7  )
8
9
10 # Decorator creates a Celery task from a regular function
11 @app.task
12 def add(x, y):
13     return x + y
```

# MINIMAL EXAMPLE – CALLING TASKS

```
1  from celery_example import add
2
3  # Fire up the task
4  add.delay(1, 4)
5
6  # Or a more sophisticated way
7  res = add.apply_async(args=(2, 4), queue='celery')
8
9  res.status    # Get status of the task
10 res.get()     # Wait for the result
```

# PROBLEM SPECIFICATION

- ▶ We have an online payment method where each order must go through non-trivial scoring process.
- ▶ Problems with synchronous code:
  - ▶ Scoring may take up to a minute.
  - ▶ The computation is resource-heavy and must not affect processing of new orders.
  - ▶ To increase throughput of the app, different scoring tasks must be run concurrently.
  - ▶ Scoring cannot run in parallel for one customer.

# IMPLEMENTATION IN CELERY

# TASK ORCHESTRATION